

# PySpark Practical Exam Paper

---

## Advanced Transformations and Actions

---

Duration: 90 minutes

Total Questions: 10

Difficulty: Intermediate to Advanced

Marks:  $10 \times 10 = 100$  marks

---

## QUESTION 1: Load and Transform Data (CSV Loading + Column Operations)

---

Marks: 10

### Question

---

You are given a CSV file `employees.csv` with the following sample data:

```
emp_id,emp_name,department,salary,joining_date
1,Alice,IT,75000,2022-01-15
2,Bob,HR,65000,2021-03-20
3,Clara,IT,85000,2020-06-10
4,David,Finance,90000,2019-11-05
5,Eve,IT,78000,2023-02-28
```

### Task:

1. Load the CSV file with header and infer schema.
2. Rename the column `emp_name` to `employee_name`.
3. Add a new column `bonus_amount` calculated as 10% of salary.
4. Add another column `department_upper` that converts the department name

to uppercase.

5. Select and display only: emp\_id, employee\_name, department\_upper, salary, bonus\_amount.

## Answer

---

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, lit, upper

# Create SparkSession
spark = SparkSession.builder.appName("EmployeeTransform").getOrCreate()

# Load CSV with header and schema inference
df = spark.read.option("header", True).option("inferSchema", True).csv("path/to/employee.csv")

# Rename column
df = df.withColumnRenamed("emp_name", "employee_name")

# Add bonus column (10% of salary)
df = df.withColumn("bonus_amount", col("salary") * 0.10)

# Add department uppercase column
df = df.withColumn("department_upper", upper(col("department")))

# Select required columns and display
result_df = df.select("emp_id", "employee_name", "department_upper", "salary", "bonus_amount")
result_df.show(truncate=False)
```

## Explanation

---

- **spark.read.option("header", True)**: Treats the first row as column names instead of data.
- **inferSchema, True**: Automatically detects data types (emp\_id as integer, salary as double).
- **withColumnRenamed()**: Changes column name from emp\_name to employee\_name for clarity.

- `col("salary") * 0.10`: References the salary column and multiplies by 0.10 (10%) to calculate bonus.
- `upper(col("department"))`: Converts text to uppercase (IT, HR, FINANCE, etc.).
- `select()`: Transformation that picks specific columns, reducing data footprint.
- `show()`: Action that triggers computation and displays results.

## Key Concepts Tested

---

- CSV file reading with options
  - Column renaming
  - Column creation with arithmetic operations
  - String functions (upper)
  - Column selection
- 

## QUESTION 2: Handling Null Values and Data Cleansing

---

Marks: 10

### Question

---

You have a customer dataset with missing values:

```
customer_id,customer_name,email,phone,city,age
101,Alice,alice@email.com,9876543210,Mumbai,28
102,Bob,,9876543211,Delhi,NULL
103,Clara,clara@email.com,,Bangalore,35
104,,david@email.com,9876543213,Chennai,42
105,Eve,eve@email.com,9876543214,NULL,29
```

Task:

1. Load the data and show the null values.
2. Remove rows where `customer_name` is NULL.
3. Fill NULL values in the phone column with "Unknown".
4. Fill NULL values in the city column with "Not Specified".
5. Drop the `email` column entirely.
6. Display the cleaned data.

## Answer

---

```

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("CustomerCleansing").getOrCreate()

# Create DataFrame with sample data
data = [
    (101, "Alice", "alice@email.com", "9876543210", "Mumbai", 28),
    (102, "Bob", None, "9876543211", "Delhi", None),
    (103, "Clara", "clara@email.com", None, "Bangalore", 35),
    (104, None, "david@email.com", "9876543213", "Chennai", 42),
    (105, "Eve", "eve@email.com", "9876543214", None, 29)
]
schema = ["customer_id", "customer_name", "email", "phone", "city", "age"]
df = spark.createDataFrame(data, schema)

# Show null values (all data for inspection)
print("Original Data with Nulls:")
df.show(truncate=False)

# Remove rows where customer_name is NULL
df_cleaned = df.dropna(subset=["customer_name"])

# Fill NULL in phone with "Unknown"
df_cleaned = df_cleaned.fillna({"phone": "Unknown", "city": "Not Specified"})

# Drop email column
df_cleaned = df_cleaned.drop("email")

# Display cleaned data
print("\nCleaned Data:")
df_cleaned.show(truncate=False)

```

## Explanation

---

- `dropna(subset=["customer_name"])`: Removes entire rows where `customer_name` is `NULL`, ensuring no invalid customer records remain.
- `fillna({...})`: Dictionary-based fill replaces `NULL` values in specific columns:

- phone → "Unknown" (indicating missing contact)
- city → "Not Specified" (placeholder for unknown location)
- **drop("email")**: Removes the email column entirely, freeing memory if not needed.
- **show()**: Action that executes the transformation chain and displays final result.

## Key Concepts Tested

---

- Null value detection
- Dropping rows with nulls in specific columns
- Filling nulls with default values
- Column dropping
- Data cleansing workflows

---

## QUESTION 3: Total Purchases by Customer (GroupBy + Aggregation)

---

Marks: 10

### Question

---

You have a sales transaction dataset:

```
transaction_id, customer_id, customer_name, product, amount, purchase_date
1,101, Alice, Laptop, 50000, 2025-01-10
2,102, Bob, Mobile, 25000, 2025-01-11
3,101, Alice, Keyboard, 5000, 2025-01-12
4,103, Clara, Mouse, 1500, 2025-01-13
5,101, Alice, Monitor, 15000, 2025-01-14
6,102, Bob, Keyboard, 5000, 2025-01-15
```

Task:

1. Load the transaction data.
2. Group by `customer_id` and `customer_name`.
3. Calculate:
  - Total purchase amount per customer
  - Number of transactions per customer
  - Average transaction amount per customer
4. Sort by total purchase amount in descending order.
5. Display the results.

## **Answer**

---

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import sum, count, avg, col

spark = SparkSession.builder.appName("SalesAnalysis").getOrCreate()

# Create sample data
data = [
    (1, 101, "Alice", "Laptop", 50000, "2025-01-10"),
    (2, 102, "Bob", "Mobile", 25000, "2025-01-11"),
    (3, 101, "Alice", "Keyboard", 5000, "2025-01-12"),
    (4, 103, "Clara", "Mouse", 1500, "2025-01-13"),
    (5, 101, "Alice", "Monitor", 15000, "2025-01-14"),
    (6, 102, "Bob", "Keyboard", 5000, "2025-01-15")
]
schema = ["transaction_id", "customer_id", "customer_name", "product", "amount", "purchase_date"]
df = spark.createDataFrame(data, schema)

# Group by customer and aggregate
customer_summary = df.groupBy("customer_id", "customer_name").agg(
    sum("amount").alias("total_purchase"),
    count("transaction_id").alias("transaction_count"),
    avg("amount").alias("avg_transaction_amount")
)
# Sort by total purchase descending
customer_summary = customer_summary.orderBy(col("total_purchase").desc())
# Display results
customer_summary.show(truncate=False)

```

## Expected Output

---

customer_id	customer_name	total_purchase	transaction_count	avg_transaction_amount
101	Alice	70000	3	23333.33
102	Bob	30000	2	15000.0
103	Clara	1500	1	1500.0

## Explanation

---

- `groupBy("customer_id", "customer_name")`: Partitions data by customer; subsequent aggregations apply within each group.
- `sum("amount").alias("total_purchase")`: Adds all amounts for each customer; `.alias()` renames the column for readability.
- `count("transaction_id").alias("transaction_count")`: Counts rows (transactions) per customer group.
- `avg("amount").alias("avg_transaction_amount")`: Computes mean transaction value per customer.
- `orderBy(col("total_purchase").desc())`: Sorts descending (highest totals first) for easy identification of top customers.

## Key Concepts Tested

---

- DataFrame creation and loading
- GroupBy operation (transformation)
- Multiple aggregations in a single `agg()` call
- Sorting results
- Aliasing columns for clarity

---

## QUESTION 4: Running Payroll Calculation (Window Functions + When/Otherwise)

---

Marks: 10

## Question

---

You have an employee payroll dataset:

```
emp_id,emp_name,department,salary,performance_rating
1,Alice,IT,75000,4.5
2,Bob,HR,65000,3.8
3,Clara,IT,85000,4.2
4,David,Finance,90000,4.8
5,Eve,IT,78000,4.0
```

## Task:

1. Load the data.
2. Add a bonus column based on performance rating:
  - Rating  $\geq 4.5$ : 15% bonus
  - Rating 4.0-4.49: 10% bonus
  - Rating < 4.0: 5% bonus
3. Calculate total compensation (salary + bonus).
4. Rank employees within each department by salary (descending).
5. Display: emp\_id, emp\_name, department, salary, bonus, total\_compensation, dept\_rank.

## Answer

---

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import col, when, round, rank
from pyspark.sql.window import Window

spark = SparkSession.builder.appName("PayrollCalculation").getOrCreate()

# Sample payroll data
data = [
    (1, "Alice", "IT", 75000, 4.5),
    (2, "Bob", "HR", 65000, 3.8),
    (3, "Clara", "IT", 85000, 4.2),
    (4, "David", "Finance", 90000, 4.8),
    (5, "Eve", "IT", 78000, 4.0)
]

schema = ["emp_id", "emp_name", "department", "salary", "performance_rating"]
df = spark.createDataFrame(data, schema)

# Add bonus column with nested when/otherwise logic
df = df.withColumn(
    "bonus",
    when(col("performance_rating") >= 4.5, col("salary") * 0.15)
        .when((col("performance_rating") >= 4.0) & (col("performance_rating") < 4.5), col("salary") * 0.1)
        .otherwise(col("salary") * 0.05)
)

# Calculate total compensation
df = df.withColumn("total_compensation", col("salary") + col("bonus"))

# Define window for department ranking
dept_window = Window.partitionBy("department").orderBy(col("salary").desc())

# Add department rank
df = df.withColumn("dept_rank", rank().over(dept_window))

# Select and display required columns
result = df.select("emp_id", "emp_name", "department", "salary", "bonus", "total_compensation")
result.show(truncate=False)

```

## Expected Output

emp_id	emp_name	department	salary	bonus	total_compensation	dept_rank
1	Alice	IT	75000	7500.0	82500.0	2
3	Clara	IT	85000	8500.0	93500.0	1
5	Eve	IT	78000	7800.0	85800.0	3
2	Bob	HR	65000	3250.0	68250.0	1
4	David	Finance	90000	13500.0	103500.0	1

## Explanation

- `when().when().otherwise()`: Nested conditional logic creates tiered bonuses:
  - Rating  $\geq 4.5 \rightarrow 15\%$  bonus
  - Rating  $[4.0, 4.5) \rightarrow 10\%$  bonus
  - Rating  $< 4.0 \rightarrow 5\%$  bonus
- `col("salary") * 0.15`: Multiplies salary by percentage; bonus is calculated relative to base pay.
- `withColumn("total_compensation", ...)`: Sums salary and bonus for gross compensation.
- `Window.partitionBy("department").orderBy(col("salary").desc())`: Creates ranking scope:
  - `partitionBy()` groups employees by department
  - `orderBy(col("salary").desc())` orders within group by salary (highest first)
- `rank().over(dept_window)`: Assigns rank; employees with same salary get same rank, next rank skips.

## Key Concepts Tested

- When/Otherwise conditional logic (nested)
- Arithmetic operations and column creation
- Window functions (`partitionBy`, `orderBy`)
- Ranking within groups
- Complex transformation chaining

## **QUESTION 5: Food and Beverage Sales Summary (GroupBy with Multiple Aggregations)**

---

**Marks:** 10

### **Question**

---

You have a product sales dataset:

```
product_id,product_name,category,quantity_sold,unit_price,stock_on_hand
1,Coffee Beans,Beverage,150,500,80
2,Tea Leaves,Beverage,200,300,120
3,Bread,Food,300,100,50
4,Milk,Beverage,250,150,40
5,Cheese,Food,100,800,30
6,Orange Juice,Beverage,180,200,60
```

### **Task:**

1. Load the product sales data.
2. Calculate:
  - Total quantity sold per category
  - Total revenue per category ( $\text{quantity\_sold} \times \text{unit\_price}$ )
  - Total stock on hand per category
  - Average unit price per category
3. Sort by total revenue in descending order.
4. Display results.

### **Answer**

---

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import sum, avg, col, round

spark = SparkSession.builder.appName("SalesAnalysis").getOrCreate()

# Sample product data
data = [
    (1, "Coffee Beans", "Beverage", 150, 500, 80),
    (2, "Tea Leaves", "Beverage", 200, 300, 120),
    (3, "Bread", "Food", 300, 100, 50),
    (4, "Milk", "Beverage", 250, 150, 40),
    (5, "Cheese", "Food", 100, 800, 30),
    (6, "Orange Juice", "Beverage", 180, 200, 60)
]

schema = ["product_id", "product_name", "category", "quantity_sold", "unit_price", "stock_on_hand"]
df = spark.createDataFrame(data, schema)

# Add revenue column
df = df.withColumn("revenue", col("quantity_sold") * col("unit_price"))

# Group by category and aggregate
category_summary = df.groupBy("category").agg(
    sum("quantity_sold").alias("total_quantity"),
    sum("revenue").alias("total_revenue"),
    sum("stock_on_hand").alias("total_stock"),
    round(avg("unit_price"), 2).alias("avg_unit_price")
)

# Sort by revenue descending
category_summary = category_summary.orderBy(col("total_revenue").desc())

# Display
category_summary.show(truncate=False)

```

## Expected Output

category	total_quantity	total_revenue	total_stock	avg_unit_price
Beverage	780	237000	320	287.5
Food	400	110000	80	450.0

## Explanation

---

- `withColumn("revenue", col("quantity_sold") * col("unit_price"))`: Creates derived column by multiplying quantity and unit price for each product.
- `groupBy("category")`: Groups all products by their category (Beverage, Food).
- `sum("quantity_sold").alias("total_quantity")`: Sums quantities within each category.
- `sum("revenue").alias("total_revenue")`: Calculates total sales revenue per category.
- `round(avg("unit_price"), 2)`: Computes average price per category, rounded to 2 decimals.
- `orderBy(col("total_revenue").desc())`: Sorts categories by revenue (highest first).

## Key Concepts Tested

---

- Derived column creation (revenue calculation)
- Multiple aggregations with aliases
- Rounding functions
- GroupBy with multiple functions
- Sorting aggregated results

---

## QUESTION 6: Discount Calculation and Pricing (Conditional Logic + Arithmetic)

---

Marks: 10

## Question

---

You have a product inventory:

```
product_id,product_name,category,original_price,quantity_in_stock
1,Laptop,Electronics,100000,15
2,Mouse,Electronics,2000,50
3,USB Cable,Electronics,500,100
4,Monitor,Electronics,30000,8
5,Keyboard,Electronics,8000,25
```

## Task:

1. Load the inventory data.
2. Apply discount logic:
  - Electronics with price > 50000: 15% discount
  - Electronics with price 10000-50000: 10% discount
  - Electronics with price < 10000: 5% discount
3. Calculate discounted price.
4. Calculate total inventory value (discounted\_price × quantity).
5. Display: product\_id, product\_name, original\_price, discount\_percent, discounted\_price, total\_inventory\_value.

## Answer

---

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, when, round

spark = SparkSession.builder.appName("DiscountCalculation").getOrCreate()

# Sample inventory data
data = [
    (1, "Laptop", "Electronics", 100000, 15),
    (2, "Mouse", "Electronics", 2000, 50),
    (3, "USB Cable", "Electronics", 500, 100),
    (4, "Monitor", "Electronics", 30000, 8),
    (5, "Keyboard", "Electronics", 8000, 25)
```

```

(4, "Monitor", "Electronics", 50000, 8),
(5, "Keyboard", "Electronics", 8000, 25)
]

schema = ["product_id", "product_name", "category", "original_price", "quantity_in_stock"]
df = spark.createDataFrame(data, schema)

# Apply tiered discount logic
df = df.withColumn(
    "discount_percent",
    when(col("original_price") > 50000, 15)
    .when((col("original_price") >= 10000) & (col("original_price") <= 50000), 10)
    .otherwise(5)
)

# Calculate discounted price
df = df.withColumn(
    "discounted_price",
    col("original_price") * (1 - col("discount_percent") / 100)
)

# Calculate total inventory value
df = df.withColumn(
    "total_inventory_value",
    col("discounted_price") * col("quantity_in_stock")
)

# Select and display required columns
result = df.select(
    "product_id",
    "product_name",
    "original_price",
    "discount_percent",
    round("discounted_price", 2).alias("discounted_price"),
    round("total_inventory_value", 2).alias("total_inventory_value")
)

result.show(truncate=False)

```

## Expected Output

---

product_id	product_name	original_price	discount_percent	discounted_price	total_inventory
1	Laptop	100000	15	85000.0	1275000.0
2	Mouse	2000	5	1900.0	95000.0
3	USB Cable	500	5	475.0	47500.0
4	Monitor	30000	10	27000.0	216000.0
5	Keyboard	8000	5	7600.0	190000.0

## Explanation

---

- `when(...).when(...).otherwise(...)`: Tiered discount logic based on price ranges.
- `col("original_price") * (1 - col("discount_percent") / 100)`: Formula to calculate discounted price:
  - Discount of 15% → multiply by 0.85
  - Discount of 10% → multiply by 0.90
  - Discount of 5% → multiply by 0.95
- `col("discounted_price") * col("quantity_in_stock")`: Inventory value based on discounted price.
- `round(..., 2)`: Rounds to 2 decimals for currency display.

## Key Concepts Tested

---

- Tiered conditional logic (when/otherwise)
- Percentage calculations
- Multiple derived columns
- Column chaining
- Rounding for financial data

## QUESTION 7: Load and Flatten Nested JSON Data (JSON Reading + Array/Explode Functions)

---

Marks: 10

## Question

---

You have a nested JSON file employees.json:

```
[  
  {  
    "emp_id": 1,  
    "emp_name": "Alice",  
    "department": "IT",  
    "skills": ["Python", "Spark", "SQL"]  
  },  
  {  
    "emp_id": 2,  
    "emp_name": "Bob",  
    "department": "HR",  
    "skills": ["Recruitment", "Payroll"]  
  },  
  {  
    "emp_id": 3,  
    "emp_name": "Clara",  
    "department": "IT",  
    "skills": ["Java", "Docker", "Kubernetes", "Python"]  
  }  
]
```

## Task:

1. Load the multiline JSON file.
2. Display the raw nested structure.
3. Flatten the data by exploding the skills array into individual rows.
4. Display flattened data.
5. Count unique skills across all employees.

## Answer

---

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import explode, col, countDistinct

spark = SparkSession.builder.appName("JSONFlattening").getOrCreate()

# Load multiline JSON file
df = spark.read.option("multiline", True).json("path/to/employees.json")

print("Original Nested Structure:")
df.show(truncate=False)
df.printSchema()

# Explode the skills array to flatten data
flattened_df = df.select(
    col("emp_id"),
    col("emp_name"),
    col("department"),
    explode(col("skills")).alias("skill")
)

print("\nFlattened Data:")
flattened_df.show(truncate=False)

# Count distinct skills
skill_count = flattened_df.select(countDistinct("skill").alias("unique_skills")).collect()
print(f"\nTotal Unique Skills: {skill_count[0][0]}")

# Alternative: Show all distinct skills
flattened_df.select("skill").distinct().show()

```

## Expected Output

---

Original:

emp_id	emp_name	department	skills
1	Alice	IT	[Python, Spark, SQL]
2	Bob	HR	[Recruitment, Payroll]
3	Clara	IT	[Java, Docker, Kubernetes, Python]

**Flattened:**

emp_id	emp_name	department	skill
1	Alice	IT	Python
1	Alice	IT	Spark
1	Alice	IT	SQL
2	Bob	HR	Recruitment
2	Bob	HR	Payroll
3	Clara	IT	Java
3	Clara	IT	Docker
3	Clara	IT	Kubernetes
3	Clara	IT	Python

## Explanation

---

- `spark.read.option("multiline", True).json()`: Reads pretty-printed JSON where arrays/objects span multiple lines.
- `explode(col("skills"))`: Expands array column into multiple rows:
  - 1 row with 3 skills → 3 rows
  - Preserves other columns (emp\_id, emp\_name, department)
- `.alias("skill")`: Renames the exploded column from "skills" to "skill" (singular).
- `countDistinct("skill")`: Counts unique skill values across all rows.
- `distinct()`: Returns unique skill values.

## Key Concepts Tested

---

- JSON file reading with multiline option
  - Schema inspection (printSchema)
  - Array explosion (row expansion)
  - Column aliasing
  - Distinct value counting
- 

## **QUESTION 8: Employee Earnings Above Department Average (Window Functions + Self-Referencing)**

---

**Marks:** 10

### **Question**

---

You have an employee salary dataset:

```
emp_id,emp_name,department,salary
1,Alice,IT,75000
2,Bob,IT,80000
3,Clara,IT,70000
4,David,Finance,90000
5,Eve,Finance,85000
6,Frank,HR,60000
```

### **Task:**

1. Load the employee data.
2. Calculate the average salary per department using a window function.
3. Identify employees whose salary is above their department's average.
4. Display: emp\_id, emp\_name, department, salary, dept\_avg\_salary, is\_above\_avg (Yes/No).
5. Sort by department and salary descending.

### **Answer**

---

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, avg, when, round
from pyspark.sql.window import Window

spark = SparkSession.builder.appName("EmployeeAnalysis").getOrCreate()

# Sample employee data
data = [
    (1, "Alice", "IT", 75000),
    (2, "Bob", "IT", 80000),
    (3, "Clara", "IT", 70000),
    (4, "David", "Finance", 90000),
    (5, "Eve", "Finance", 85000),
    (6, "Frank", "HR", 60000)
]

schema = ["emp_id", "emp_name", "department", "salary"]
df = spark.createDataFrame(data, schema)

# Define window specification for department average
dept_window = Window.partitionBy("department")

# Add department average salary column using window function
df = df.withColumn(
    "dept_avg_salary",
    round(avg(col("salary")).over(dept_window), 2)
)

# Determine if employee salary is above department average
df = df.withColumn(
    "is_above_avg",
    when(col("salary") > col("dept_avg_salary"), "Yes").otherwise("No")
)

# Sort by department and salary (descending)
result = df.orderBy(col("department"), col("salary").desc())

# Display required columns
result.select("emp_id", "emp_name", "department", "salary", "dept_avg_salary", "is_above_avg").show()
```

## Expected Output

---

emp_id	emp_name	department	salary	dept_avg_salary	is_above_avg
2	Bob	Finance	90000	87500.0	Yes
5	Eve	Finance	85000	87500.0	No
6	Frank	HR	60000	60000.0	No
2	Bob	IT	80000	75000.0	Yes
1	Alice	IT	75000	75000.0	No
3	Clara	IT	70000	75000.0	No

## Explanation

---

- `Window.partitionBy("department")`: Creates window scope per department. All employees in IT see only IT average; Finance employees see Finance average.
- `avg(col("salary")).over(dept_window)`: Computes average salary within each partition without grouping. All rows from a department retain their original records plus the department average.
- `when(col("salary") > col("dept_avg_salary"), "Yes").otherwise("No")`: Compares individual salary to department average:
  - If salary > avg → "Yes"
  - Else → "No"
- `orderBy(col("department"), col("salary").desc())`: Sorts by department first, then salary descending within each department.

## Key Concepts Tested

---

- Window functions with `partitionBy`
- Aggregate functions within windows (`avg`)
- Conditional comparisons
- Multi-column sorting
- Self-referencing columns in window context

## QUESTION 9: Remove Duplicates Using Window Functions

---

Marks: 10

### Question

---

You have a customer dataset with duplicates:

```
customer_id, customer_name, email, city, signup_date
101, Alice, alice@email.com, Mumbai, 2023-01-15
102, Bob, bob@email.com, Delhi, 2023-02-20
101, Alice, alice@email.com, Mumbai, 2023-01-15
103, Clara, clara@email.com, Bangalore, 2023-03-10
102, Bob, bob@email.com, Delhi, 2023-02-20
104, David, david@email.com, Chennai, 2023-04-05
```

### Task:

1. Load the customer data with duplicates.
2. Show the original data with duplicate rows.
3. Use a window function to assign a row number within each customer group (ordered by `signup_date`).
4. Keep only the first occurrence of each customer (`row_number = 1`).
5. Display the deduplicated data.

### Answer

---

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import row_number
from pyspark.sql.window import Window

spark = SparkSession.builder.appName("DeduplicationWindow").getOrCreate()

# Sample data with duplicates
data = [
    (101, "Alice", "alice@email.com", "Mumbai", "2023-01-15"),
    (102, "Bob", "bob@email.com", "Delhi", "2023-02-20"),
    (101, "Alice", "alice@email.com", "Mumbai", "2023-01-15"),
    (103, "Clara", "clara@email.com", "Bangalore", "2023-03-10"),
    (102, "Bob", "bob@email.com", "Delhi", "2023-02-20"),
    (104, "David", "david@email.com", "Chennai", "2023-04-05")
]

schema = ["customer_id", "customer_name", "email", "city", "signup_date"]
df = spark.createDataFrame(data, schema)

print("Original Data with Duplicates:")
df.show(truncate=False)
print(f"Total Records: {df.count()}")

# Define window to assign row number per customer (ordered by signup_date)
window_spec = Window.partitionBy("customer_id").orderBy("signup_date")

# Add row number column
df_with_rn = df.withColumn("row_num", row_number().over(window_spec))

# Keep only row_number = 1 (first occurrence)
df_deduplicated = df_with_rn.filter("row_num = 1").drop("row_num")

print("\nDeduplicated Data:")
df_deduplicated.show(truncate=False)
print(f"Total Unique Records: {df_deduplicated.count()}")

```

## Expected Output

---

Before:

customer_id	customer_name	email	city	signup_date
101	Alice	alice@email.com	Mumbai	2023-01-15
102	Bob	bob@email.com	Delhi	2023-02-20
101	Alice	alice@email.com	Mumbai	2023-01-15
103	Clara	clara@email.com	Bangalore	2023-03-10
102	Bob	bob@email.com	Delhi	2023-02-20
104	David	david@email.com	Chennai	2023-04-05

Total Records: 6

**After:**

customer_id	customer_name	email	city	signup_date
101	Alice	alice@email.com	Mumbai	2023-01-15
102	Bob	bob@email.com	Delhi	2023-02-20
103	Clara	clara@email.com	Bangalore	2023-03-10
104	David	david@email.com	Chennai	2023-04-05

Total Unique Records: 4

## Explanation

- `Window.partitionBy("customer_id").orderBy("signup_date")`: Creates a window for each customer, ordered by signup date (earliest first).
- `row_number().over(window_spec)`: Assigns sequential numbers (1, 2, 3...) within each customer group:
  - First occurrence of customer 101 → `row_num = 1`
  - Duplicate of customer 101 → `row_num = 2`
- `filter("row_num = 1")`: Keeps only the first occurrence of each customer.
- `drop("row_num")`: Removes the temporary `row_number` column.

## Key Concepts Tested

- Window functions for deduplication

- Row numbering within partitions
  - Filtering based on row numbers
  - Preserving first occurrences
  - Distinct counting
- 

## **QUESTION 10: Word Count Program in PySpark (RDD Transformations + Actions)**

---

**Marks:** 10

### **Question**

---

You have a text file sample.txt containing:

```
Hello world hello spark  
Spark is big data Spark  
Hello hello Spark world
```

### **Task:**

1. Load the text file using RDD.
2. Split each line into words.
3. Convert all words to lowercase.
4. Create (word, 1) pairs for each word.
5. Aggregate counts for each word.
6. Sort by count descending.
7. Display top 10 words and their counts.

### **Answer**

---

```

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("WordCount").getOrCreate()
sc = spark.sparkContext

# Load text file as RDD
text_rdd = sc.textFile("path/to/sample.txt")

# Perform word count transformations
word_counts = (text_rdd
    .flatMap(lambda line: line.split())          # Split each line into words
    .map(lambda word: word.lower())              # Convert to lowercase
    .map(lambda word: (word, 1))                 # Create (word, 1) pairs
    .reduceByKey(lambda a, b: a + b)             # Sum counts for each word
    .sortBy(lambda x: x[1], ascending=False)      # Sort by count descending
)

# Action: Collect and display top 10
top_words = word_counts.take(10)

print("Word Count Results (Top 10):")
for word, count in top_words:
    print(f"{word}: {count}")

# Alternative: Save to file
word_counts.saveAsTextFile("path/to/output/wordcount")

# Show total unique words
total_unique = word_counts.count()
print(f"\nTotal Unique Words: {total_unique}")

```

## Expected Output

---

Word Count Results (Top 10):

hello: 4

spark: 4

world: 2

is: 1

big: 1

data: 1

Total Unique Words: 6

## Explanation

---

- **sc.textFile()**: Reads text file and returns RDD where each element is a line.
- **flatMap(lambda line: line.split())**: Transformation that:
  - Splits each line by whitespace into words
  - Flattens results (flatMap returns single RDD of words, not nested)
  - Example: ["Hello world", "Spark"] → ["Hello", "world", "Spark"]
- **map(lambda word: word.lower())**: Transformation converting each word to lowercase.
- **map(lambda word: (word, 1))**: Transformation creating (word, 1) tuples for counting.
- **reduceByKey(lambda a, b: a + b)**: Transformation that:
  - Groups by key (word)
  - Aggregates values using lambda function (a + b adds counts)
  - Result: (word, total\_count) pairs
- **sortBy(lambda x: x[1], ascending=False)**: Transformation sorting by count (index 1) descending.
- **take(10)**: Action that returns first 10 elements from RDD.
- **saveAsTextFile()**: Action that writes results to HDFS/file system.
- **count()**: Action that returns total number of elements.

## Key Concepts Tested

---

- RDD creation from text files

- flatMap transformation
  - map transformation (multiple chained)
  - reduceByKey aggregation
  - Sorting RDD data
  - RDD actions (take, count, saveAsTextFile)
  - Lambda functions in RDD operations
- 

## SUMMARY OF KEY CONCEPTS COVERED

Concept	Questions
DataFrame Creation & Loading	Q1, Q2, Q3, Q4, Q5
CSV/JSON File Reading	Q1, Q7
Column Operations (select, rename, add, drop)	Q1, Q6, Q8
Filtering & Conditions	Q2, Q8
When/Otherwise Logic	Q4, Q6
GroupBy & Aggregations	Q3, Q5
Window Functions	Q4, Q8, Q9
Arithmetic Operations	Q1, Q6
String Functions	Q1, Q7
Array/Explode Functions	Q7
Null Handling	Q2
Sorting/Ordering	Q3, Q4, Q8, Q10
Joins	- (Covered in extended questions)
RDD Operations	Q10

Concept	Questions
Actions (show, count, take, collect)	All questions

## SCORING RUBRIC

---

Each question (10 marks) is evaluated as follows:

- **Code Correctness (5 marks):** Code runs without errors; transformations are applied correctly.
- **Output Accuracy (3 marks):** Output matches expected results; logic is sound.
- **Explanation (2 marks):** Clear explanation of logic, transformations, and concepts.

**Pass:**  $\geq 60$  marks (6/10 questions correct)

**Merit:**  $\geq 80$  marks (8/10 questions correct)

**Distinction:**  $\geq 90$  marks (9/10 questions correct)

---

## INSTRUCTIONS FOR STUDENTS

---

### 1. Environment Setup:

- Ensure PySpark is installed and SparkSession is initialized.
- All sample data is provided in questions; use it as-is.

### 2. Execution:

- Write and test code in a Jupyter notebook or PySpark shell.
- Verify output matches expected results.
- Document your understanding of transformations and actions.

### 3. Answer Format:

- Submit Python code for each question.
- Include brief explanations of each transformation step.

- Show actual output/results.

#### 4. Time Management:

- ~9 minutes per question on average.
  - Start with easier questions (Q1, Q2, Q3) if time-pressed.
  - Advanced questions (Q7, Q8, Q9, Q10) can take more time.
- 

**Good Luck! Practice, Understand, Excel!**